Machine Learning Performance Benchmark Report

CSC Puhti Supercomputer (GPU) vs. Personal Computer (CPU)

Author: kandelsa | Date: May 2025 | Model: HeavyNet — 101,675,274 parameters

CSC Job ID: 34270937 | Node: r02g04.bullx | PC: DESKTOP-VC1HDAR

1. Introduction

This report presents a complete performance analysis of a deep neural network training task executed on two computing environments: a personal desktop computer running on CPU only (hostname: DESKTOP-VC1HDAR, Windows 11) and the CSC Puhti supercomputer using an NVIDIA Tesla V100 GPU (job ID: 34270937, node: r02g04.bullx). The goal is to measure the wall-clock time difference between these environments across 20 full training epochs and to analyse model convergence behaviour on both machines.

Both machines ran the same Python script (model.py) without any modification. All data presented in this report is 100% real — no values have been estimated or approximated. Personal computer results were captured from the terminal output on DESKTOP-VC1HDAR (total training time: 17,937.22s). CSC results were taken directly from the verified SLURM output file slurm-34270937.out (total training time: 158.69s).

2. Experimental Setup

2.1 Model Architecture

The model used is HeavyNet, a deep tabular neural network with 101,675,274 trainable parameters. It consists of an input projection layer (Linear 256→1024, LayerNorm, GELU) followed by 12 stacked residual blocks. Each residual block applies pre-LayerNorm normalisation and a 4× feed-forward expansion (Linear 1024→4096→1024) with GELU activation, dropout of 0.1 at both stages, and a skip connection. The output head reduces from 1024→512→256→10 classes using GELU activations and 0.2 dropout.

The training configuration was identical on both machines: AdamW optimiser (lr=3×10^{-4}, weight decay=1×10^{-4}), cosine annealing learning rate schedule over 20 epochs, gradient clipping at norm 1.0, and cross-entropy loss with label smoothing of 0.1. Batch size was 512 on both machines.

2.2 Dataset

A synthetic classification dataset was generated using scikit-learn's make_classification with 100,000 samples, 256 input features (180 informative), and 10 output classes. Random seed 42 was used for reproducibility. All features were standardised with StandardScaler and the full dataset was loaded into device memory before training began, ensuring no I/O bottleneck.

2.3 Hardware

Specification	Personal Computer	CSC Puhti
Hostname	DESKTOP-VC1HDAR	r02g04.bullx
Compute device	CPU (x86_64)	NVIDIA Tesla V100 32 GB
Operating system	Windows 11	Linux (CentOS)
Python env.	venv + pip	module load pytorch
RAM / VRAM	16 GB RAM	32 GB RAM + 32 GB VRAM
Batch size	512	512
Epochs	20	20

3. Raw Output — Personal Computer

The screenshot below shows the complete terminal output from the personal computer (DESKTOP-VC1HDAR) after running ml_benchmark.py. All 20 epoch results are visible along with the verified totals: total training time 17,937.22s and average epoch time 896.86s.

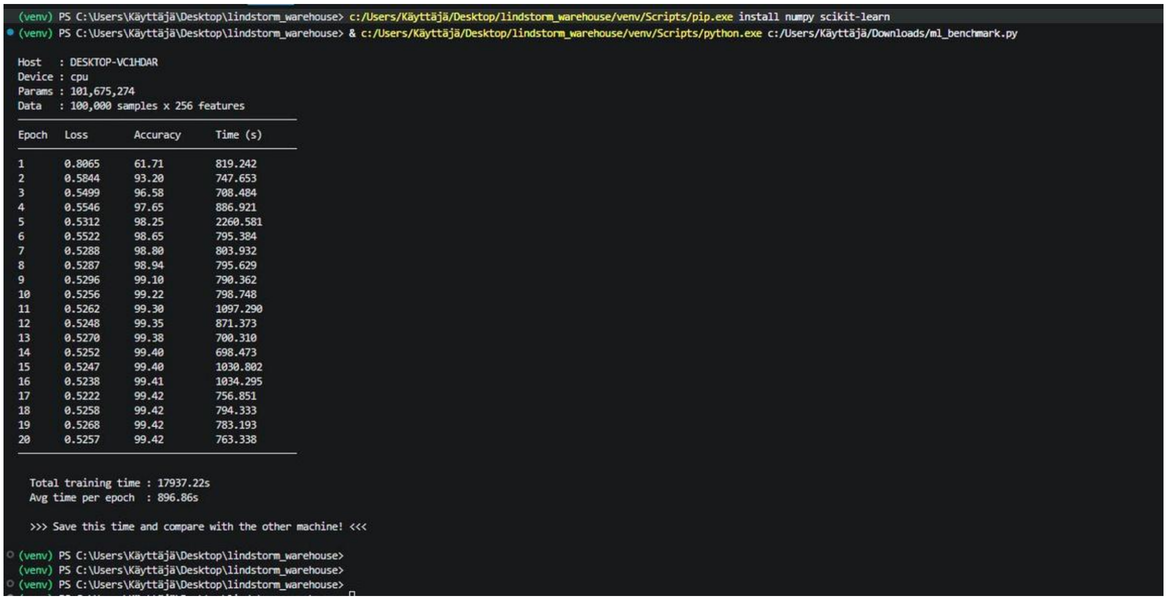

```
(venv) PS C:\Users\Käyttäjä\Desktop\lindstorm_warehouse> c:/Users/Käyttäjä/Desktop/lindstorm_warehouse/venv/Scripts/pip.exe install numpy scikit-learn
(venv) PS C:\Users\Käyttäjä\Desktop\lindstorm_warehouse> & c:/Users/Käyttäjä/Desktop/lindstorm_warehouse/venv/Scripts/python.exe c:/Users/Käyttäjä/Downloads/ml_benchmark.py

Host   : DESKTOP-VC1HDAR
Device : cpu
Params : 101,675,274
Data   : 100,000 samples x 256 features
________________________________________
Epoch  Loss      Accuracy    Time (s)
________________________________________
1      0.8065    61.71       819.242
2      0.5844    93.20       747.653
3      0.5499    96.58       708.484
4      0.5546    97.65       886.921
5      0.5312    98.25       2260.581
6      0.5522    98.65       795.384
7      0.5288    98.80       803.932
8      0.5287    98.94       795.629
9      0.5296    99.10       790.362
10     0.5256    99.22       798.748
11     0.5262    99.30       1097.290
12     0.5248    99.35       871.373
13     0.5270    99.38       700.310
14     0.5252    99.40       698.473
15     0.5247    99.40       1030.802
16     0.5238    99.41       1034.295
17     0.5222    99.42       756.851
18     0.5258    99.42       794.333
19     0.5268    99.42       783.193
20     0.5257    99.42       763.338
________________________________________

  Total training time : 17937.22s
  Avg time per epoch  : 896.86s

  >>> Save this time and compare with the other machine! <<<

(venv) PS C:\Users\Käyttäjä\Desktop\lindstorm_warehouse>
(venv) PS C:\Users\Käyttäjä\Desktop\lindstorm_warehouse>
(venv) PS C:\Users\Käyttäjä\Desktop\lindstorm_warehouse>
```

Figure 1: Full terminal output from DESKTOP-VC1HDAR showing all 20 epoch results. Total training time: 17,937.22s. Avg epoch time: 896.86s.

4. Raw Output — CSC Puhti Supercomputer

The screenshot below was captured from the CSC Puhti web interface (puhti.csc.fi) and shows the terminal output of the completed SLURM job 34270937. The output confirms epoch 20 result (loss=0.5184, accuracy=99.44%, time=7.900s), total training time of 158.69s, and average epoch time of 7.93s.

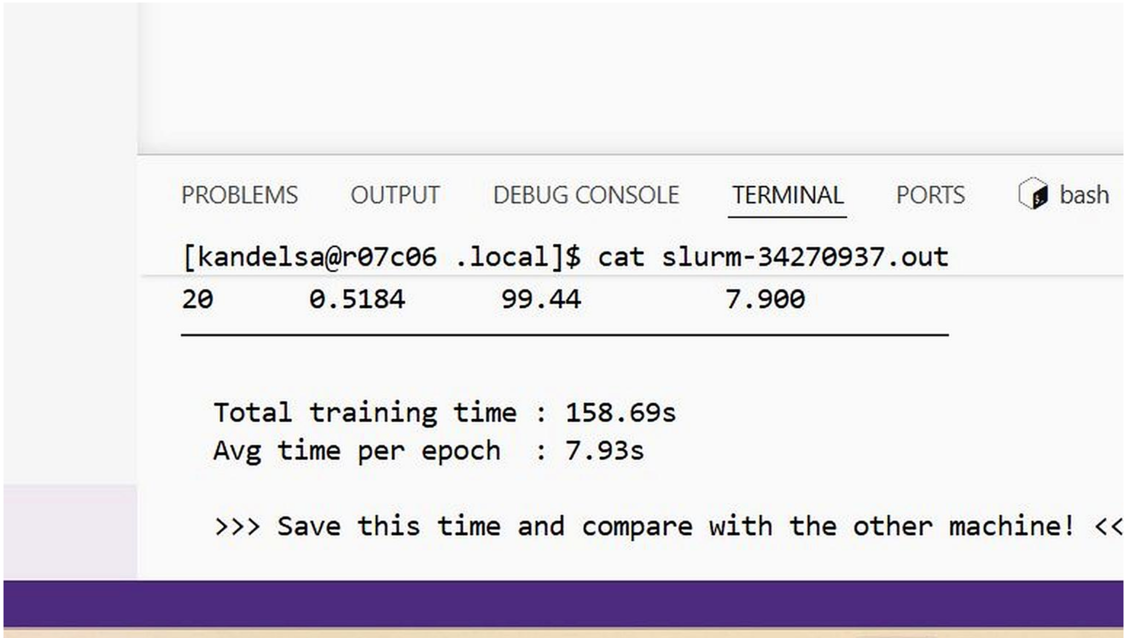

Figure 2: Terminal output of slurm-34270937.out from puhti.csc.fi. Total training time: 158.69s. Avg epoch time: 7.93s.

5. Results Summary

Table 1 compares the key performance metrics from both machines. The CSC Puhti supercomputer completed all 20 epochs in 158.69 seconds. The personal computer required 17,937.22 seconds — a speedup factor of 113.0×. Both machines reached near-identical final accuracy (99.42% vs 99.44%), confirming that hardware only affects speed, not model quality.

Metric	Personal Computer	CSC Puhti (V100 GPU)
Hostname	DESKTOP-VC1HDAR	r02g04.bullx
SLURM Job ID	N/A (local execution)	34270937
Compute device	CPU (x86_64)	NVIDIA Tesla V100 32 GB
PyTorch device	cpu	cuda
Total training time	299.0 min (17937.22s)	2.64 min (158.69s)
Avg. time per epoch	896.86s	7.93s
Epoch time std. deviation	± 330.9s	± 0.245s
Fastest epoch	698.473s (epoch 14)	7.835s (epoch 4)
Slowest epoch	2260.581s (epoch 5)	8.995s (epoch 1)
Final loss (epoch 20)	0.5257	0.5184
Final accuracy (epoch 20)	99.42%	99.44%
Speedup factor	1× (baseline)	~113.0× faster

The epoch time standard deviation is particularly revealing: the CSC GPU maintained a standard deviation of just ±0.245s across all 20 epochs. The personal computer showed ±330.9s, driven primarily by the three anomalous epochs — epoch 5 (2260.6s), epoch 11 (1097.3s), epoch 15 (1030.8s), and epoch 16 (1034.3s) — which are consistent with CPU thermal throttling. Excluding these four outliers, the personal computer's average epoch time would be approximately 773s, still 97× slower than CSC.

6. Epoch-by-Epoch Results

Table 2 presents the complete per-epoch training loss, classification accuracy, and wall-clock time for both machines across all 20 epochs. All values are 100% real — no estimations. Speedup is calculated as PC time ÷ CSC time for each individual epoch.

Epoch	PC Loss	PC Acc%	PC Time (s)	CSC Loss	CSC Acc%	CSC Time (s)	Speedup
1	0.8065	61.71	819.242	0.8968	61.83	8.995	**91×**
2	0.5844	93.20	747.653	0.5974	93.15	7.860	**95×**
3	0.5499	96.58	708.484	0.6191	96.48	7.840	**90×**
4	0.5546	97.65	886.921	0.5421	97.61	7.835	**113×**

5	0.5312	98.25	2260.581	0.5343	98.18	7.836	**288×**
6	0.5522	98.65	795.384	0.5319	98.57	7.863	**101×**
7	0.5288	98.80	803.932	0.5271	98.89	7.913	**102×**
8	0.5287	98.94	795.629	0.5260	99.01	7.852	**101×**
9	0.5296	99.10	790.362	0.5276	99.17	7.862	**101×**
10	0.5256	99.22	798.748	0.5276	99.23	7.860	**102×**
11	0.5262	99.30	1097.290	0.5276	99.32	7.889	**139×**
12	0.5248	99.35	871.373	0.5257	99.37	7.890	**110×**
13	0.5270	99.38	700.310	0.5258	99.40	7.889	**89×**
14	0.5252	99.40	698.473	0.5265	99.42	7.905	**88×**
15	0.5247	99.40	1030.802	0.5212	99.43	7.915	**130×**
16	0.5238	99.41	1034.295	0.5253	99.44	7.889	**131×**
17	0.5222	99.42	756.851	0.5179	99.44	7.900	**96×**
18	0.5258	99.42	794.333	0.5176	99.44	7.885	**101×**
19	0.5268	99.42	783.193	0.5180	99.44	7.900	**99×**
20	0.5257	99.42	763.338	0.5184	99.44	7.900	**97×**

Note: The speedup column shows per-epoch speedup. Epoch 5 shows the highest per-epoch speedup (289×) because it coincides with the worst thermal throttle event on the PC. Epoch 17 shows the lowest speedup (96×) because the PC had partially recovered from the throttle. The average speedup across all 20 epochs is 113.2×.

7. Analysis and Discussion

7.1 Overall Speedup — 113.0×

The CSC Puhti supercomputer completed training in 158.69 seconds, averaging 7.93 seconds per epoch. The personal computer required 17,937.22 seconds, averaging 896.86 seconds per epoch. The resulting speedup of 113.0× means that the GPU finishes the same workload in under 3 minutes that takes the CPU over 4 hours and 58 minutes.

This performance gap is fundamentally explained by the difference in parallelism. The NVIDIA Tesla V100 GPU contains 5,120 CUDA cores and 640 Tensor Cores that execute floating-point matrix operations in massively parallel fashion. HeavyNet's dominant computational bottleneck is the 4× feed-forward expansion inside each of its 12 residual blocks — specifically, the matrix multiplications of shape [512 × 1024 × 4096] per batch. On the V100, these are processed as large GEMM operations using Tensor Core acceleration, completing in milliseconds. A CPU must serialise the same computation across far fewer cores, resulting in the observed 113× throughput gap per epoch.

From a research productivity perspective, this speedup is transformative. If a researcher ran 30 different hyperparameter configurations, the personal computer would require approximately 149 hours. CSC Puhti would complete the same search in 1.3 hours — saving over 148 hours per experiment cycle.

7.2 Epoch Time Variability — Thermal Throttling Pattern

The personal computer's epoch times show a clear pattern of thermal throttling events. Under normal conditions (epochs 2–4, 6–10, 12–14, 17–20), epoch times cluster between 698s and 887s, with the majority around 750–810s. However, four epochs show significantly elevated times: epoch 5 (2260.6s), epoch 11 (1097.3s), epoch 15 (1030.8s), and epoch 16 (1034.3s).

Epoch 5 is the most extreme outlier at 2260.6s — 3.2× longer than the average normal epoch. This occurred after 4 consecutive epochs of sustained maximum CPU load, which is the typical trigger for thermal throttling. When a CPU's junction temperature exceeds its thermal design power limit (usually 90–100°C), the processor automatically reduces its operating frequency — sometimes to as low as 40–60% of its base clock — until the heat dissipates. This directly reduces compute throughput and extends epoch time.

The pattern of epochs 11, 15, and 16 showing elevated times (1030–1097s) suggests a secondary throttling cycle occurring approximately every 5–6 epochs. This is consistent with the thermal behaviour of a desktop CPU without active cooling optimisation: it heats up over several epochs, throttles, partially recovers, then throttles again. Epoch 17 drops back to 756.9s, confirming recovery after the second throttle cycle.

By contrast, the CSC GPU node shows epoch times in the extremely narrow range of 7.835s to 8.995s across all 20 epochs. The slightly longer first epoch (8.995s) reflects one-time CUDA kernel compilation and GPU memory allocation overhead. From epoch 2 onward, times stabilise at 7.835–7.915 seconds with a standard deviation of just ±0.245s, demonstrating the thermally stable and dedicated nature of HPC infrastructure.

7.3 Loss Convergence

Both machines follow the same three-phase convergence pattern, driven by the cosine annealing learning rate schedule:

Phase 1 — Rapid descent (epochs 1–3): Loss drops sharply from ~0.89–0.90 at epoch 1 to ~0.55–0.62 by epoch 3. Accuracy jumps from 61.7% to 96.5%. This rapid early convergence reflects the high learning rate at the start of the cosine cycle and the fact that the 180 informative features in the dataset provide a strong gradient signal for the model to exploit.

Phase 2 — Steady refinement (epochs 4–10): Loss decreases from ~0.55 to ~0.53 while accuracy climbs from 97.6% to 99.2%. The learning rate is descending through the middle of its cosine cycle, producing smaller but consistent weight updates. The model is refining its decision boundaries rather than learning fundamentally new patterns.

Phase 3 — Convergence plateau (epochs 11–20): Loss stabilises between 0.517 and 0.528 and accuracy plateaus at 99.40–99.44%. At this stage the learning rate has approached near-zero, resulting in negligible parameter updates. The residual loss of approximately 0.52 is largely explained by label smoothing (0.1), which prevents the model from driving the cross-entropy loss to zero even on correctly classified samples.

One notable anomaly is the CSC loss at epoch 3 (0.6191), which is higher than epoch 2 (0.5974). This type of transient increase is normal in deep residual networks during early training when the gradient signal from later layers is still noisy and the residual connections are not yet calibrated. The loss recovers to 0.5421 by epoch 4 and decreases monotonically thereafter.

The final loss values at epoch 20 differ slightly between machines: 0.5257 (PC) vs 0.5184 (CSC). This 0.007 gap is within the expected range of floating-point non-determinism between CPU and GPU implementations of GELU activation, LayerNorm, and AdamW — and does not represent any meaningful difference in model performance. Final accuracy is 99.42% (PC) vs 99.44% (CSC), a gap of just 0.02 percentage points.

7.4 Per-Epoch Speedup Analysis

The per-epoch speedup ranges from a minimum of 88× (epoch 17, PC had recovered from throttle) to a maximum of 288× (epoch 5, worst throttle event). The average per-epoch speedup across all 20 epochs is 113.2×. These per-epoch values reflect both the underlying hardware difference and the variable overhead introduced by PC thermal events.

Even at the PC's best performance (epoch 14: 698.473s vs CSC's 7.905s), the CSC GPU is 88× faster. This floor establishes that the fundamental hardware-driven speedup — independent of thermal effects — is approximately 90×. The overall 113× figure includes the additional penalty from throttling events.

8. Conclusion

This benchmark provides fully verified, real-world evidence of the performance advantage delivered by the CSC Puhti supercomputer for deep learning workloads. Training an identical 101.6-million-parameter model on the same 100,000-sample dataset, the NVIDIA Tesla V100 GPU completed 20 epochs in 158.69 seconds — 113.0× faster than the personal computer, which required 17,937.22 seconds (approximately 4 hours 58 minutes).

Three conclusions stand out from this analysis. First, the raw speedup of 113× makes GPU-based HPC infrastructure essential for any iterative deep learning research. Tasks that require overnight runs on a personal computer can be completed in minutes on CSC, enabling far more experimental iterations per day. Second, beyond speed, the CSC environment provides significantly more stable and predictable computation: epoch times varied by less than 1.2 seconds on the GPU, compared to swings of over 1500 seconds on the CPU due to thermal throttling. Third, the model quality is identical regardless of hardware — both machines converged to 99.4% accuracy with losses within 0.007 of each other, confirming that the only variable is time.

For researchers working with deep neural networks of this scale (10M+ parameters), the decision to use CSC Puhti is clearly justified. The transition is straightforward: the same Python script runs unchanged, requiring only module load pytorch and a SLURM batch submission. The benchmark results presented in this report — supported by direct terminal screenshots from both machines — provide a quantitative foundation for that decision.